

# Dynamic libraries and how to optimize them

C++ Russia 2024

# About me

- Yuri “yugr” Gribov
- Compiler engineer
- Gmail: tetra2005
- [t.me/the\\_real\\_yugr](https://t.me/the_real_yugr)
- <https://github.com/yugr>
- <https://www.linkedin.com/in/yugr/>



# Plan of the talk

- Dynamic libraries

# Plan of the talk

- Dynamic libraries
  - Differences from static libraries
  - Work principles
  - Pros and cons

# Plan of the talk

- Dynamic libraries
  - Differences from static libraries
  - Work principles
  - Pros and cons
- Comparison of Linux and Windows implementations

# Plan of the talk


- Dynamic libraries
  - Differences from static libraries
  - Work principles
  - Pros and cons
- Comparison of Linux and Windows implementations
- Speeding up dynamic libraries

# Plan of the talk

- Dynamic libraries
  - Differences from static libraries
  - Work principles
  - Pros and cons
- Comparison of Linux and Windows implementations
- Speeding up dynamic libraries
  - Overheads
  - And ways to reduce them

# Libraries

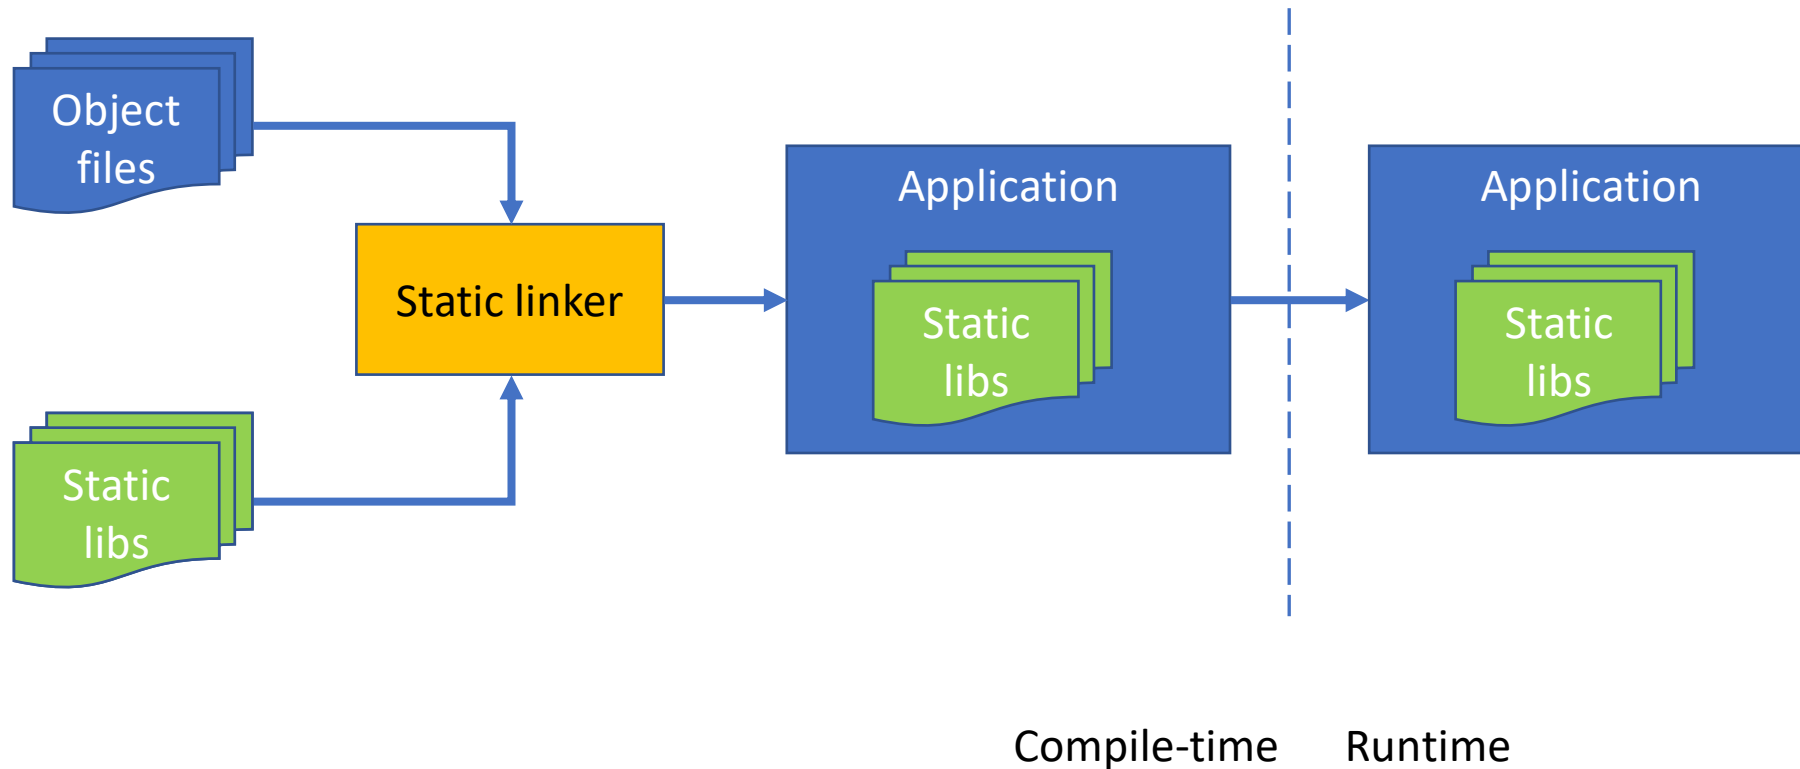
- Archives of reusable code
- Can be reused in multiple programs
- Depending on library link time can be
  - Static (.a, .lib)
  - Dynamic (.so, .dll, .dylib)
- All popular platforms support both
  - Windows, Linux, macOS, BSDs





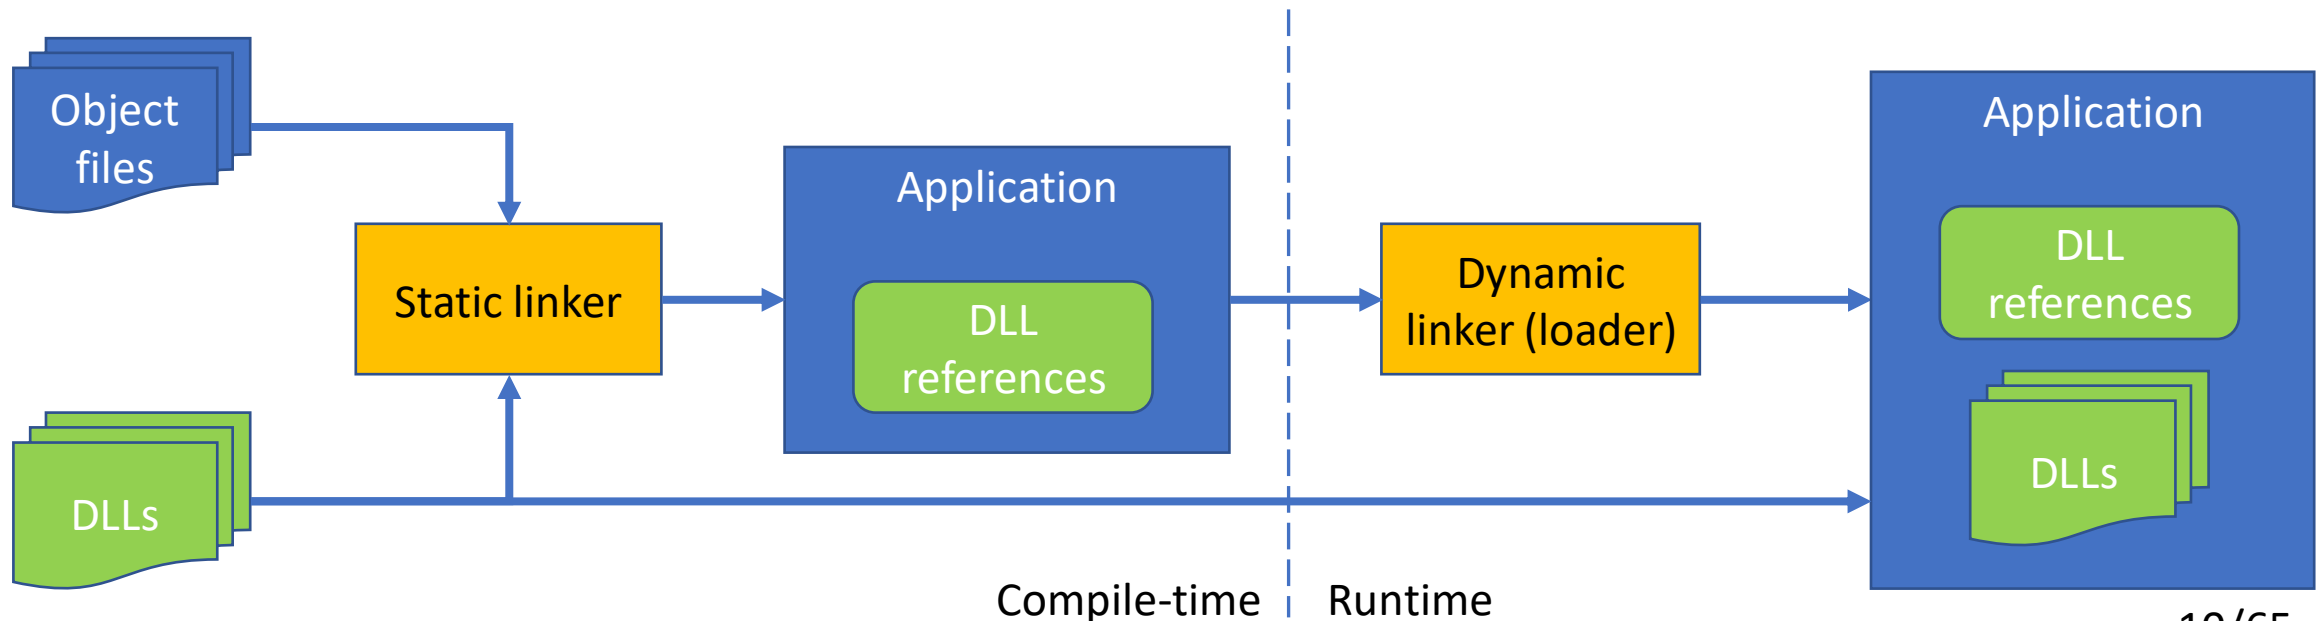
# Static libraries

- Become part of executable file at link-time



# Dynamic libraries

- Dynamic-link libraries (DLL), shared libraries, shared objects
- Not part of program executable file
- (Usually) loaded at program startup



# Using dynamic libraries

- Two main approaches:

- Traditional, link-time

```
gcc program.o -lgmp
link.exe program.obj libgmp.lib
```

- Run-time loading (dynamic loading)

```
void *lib = dlopen("libgmp.so", RTLD_LAZY | RTLD_GLOBAL);
HANDLE lib = LoadLibrary("libgmp.dll");
```

- With traditional approach library will be loaded at program startup
- With runtime loading – at any time, in any point in program
  - Enables lazy loading, plugins, etc.

# DLL advantages

- RAM and disk savings
  - ~1.1G RAM on my Ubuntu Desktop<sup>1,2</sup> (with running Firefox/KOffice/Thunderbird)
  - ~10G HDD on my Ubuntu Desktop (with Firefox, KOffice, etc.)
- Faster system updates
  - No need to recompile dependent executables on minor library updates
- Support for interesting work scenarios:
  - Lazy loading
  - Extend program functionality with user plugins
  - Load different library versions depending on environment (e.g. on processor capabilities)

1) Experiment details are available in additional slides at  
<https://github.com/yugr/CppRussia/tree/master/2024>

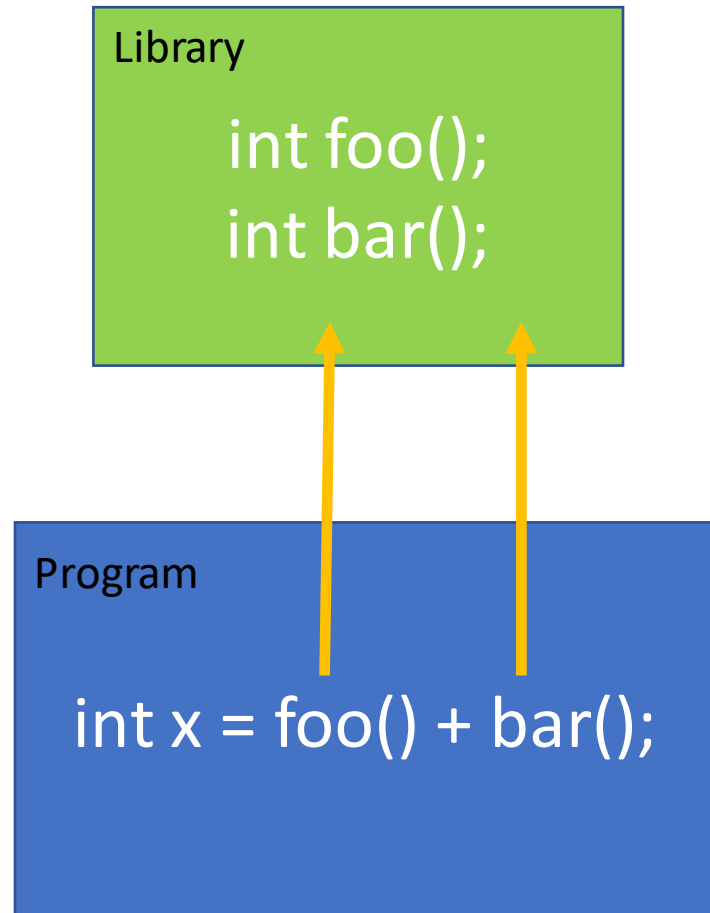
2) Based on <https://zvrba.net/articles/solib-memory-savings.html>



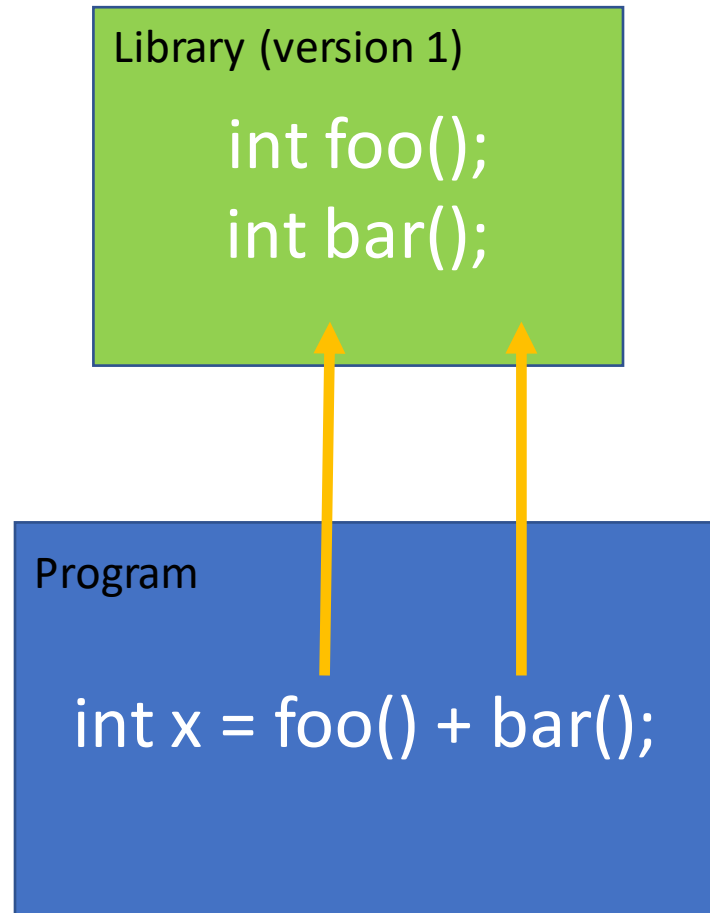
# DLL disadvantages

- Performance overhead
  - Program startup (search and load libraries, search for symbols)
  - Calling library functions
- Fragile infrastructure (DLL hell)


# DLL Hell: example



# DLL Hell: example




# DLL Hell: example





# DLL Hell: example



# DLL Hell

- It is very easy to introduce *incompatible changes* when developing a library
  - Remove function or change its signature
- Programs which used old version will not be able to work with new one
  - Will crash at startup or later

# DLL Hell: solution

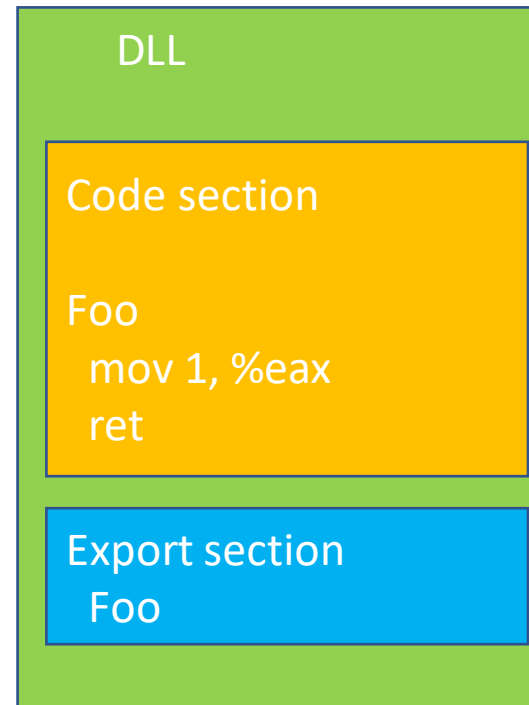
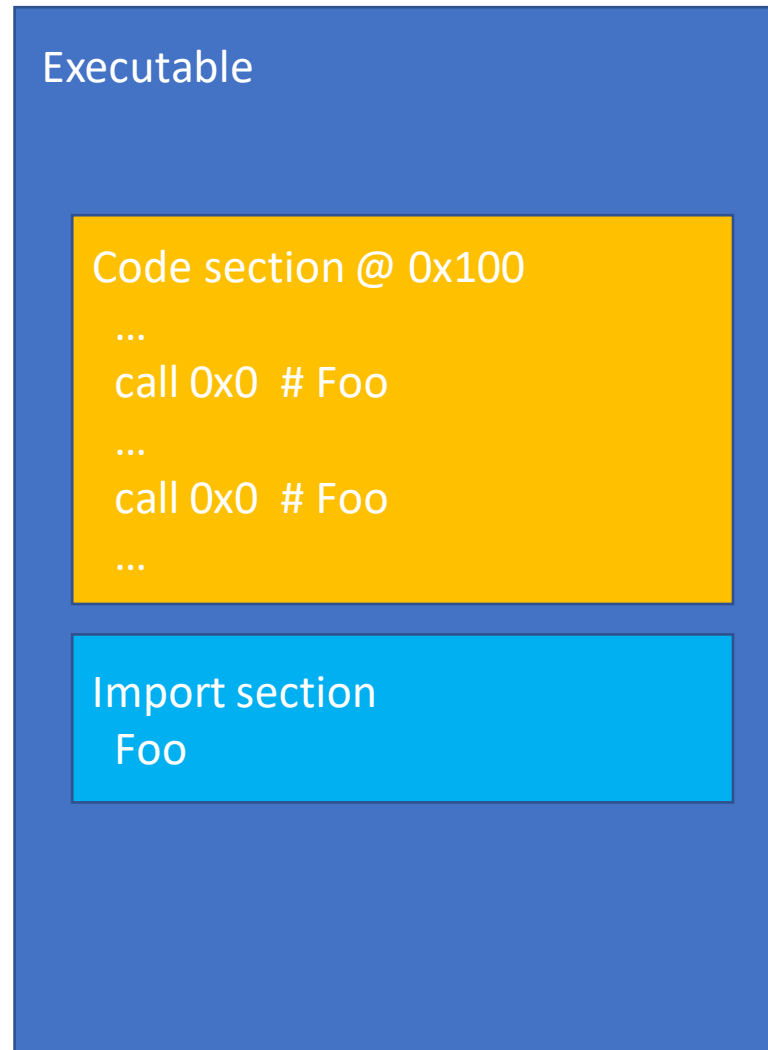
- Library developers should avoid incompatible changes
  - Incompatibility checking can and should be automated (libabigail, ABI Compliance Checker, etc.)
- If such changes are inevitable developer needs to update library version info
  - Embedded in library file
  - SONAME on Linux, DLL manifests on Windows
- This will allow OS to determine which library version is needed for particular program
- Details are OS-specific

DLL working principles

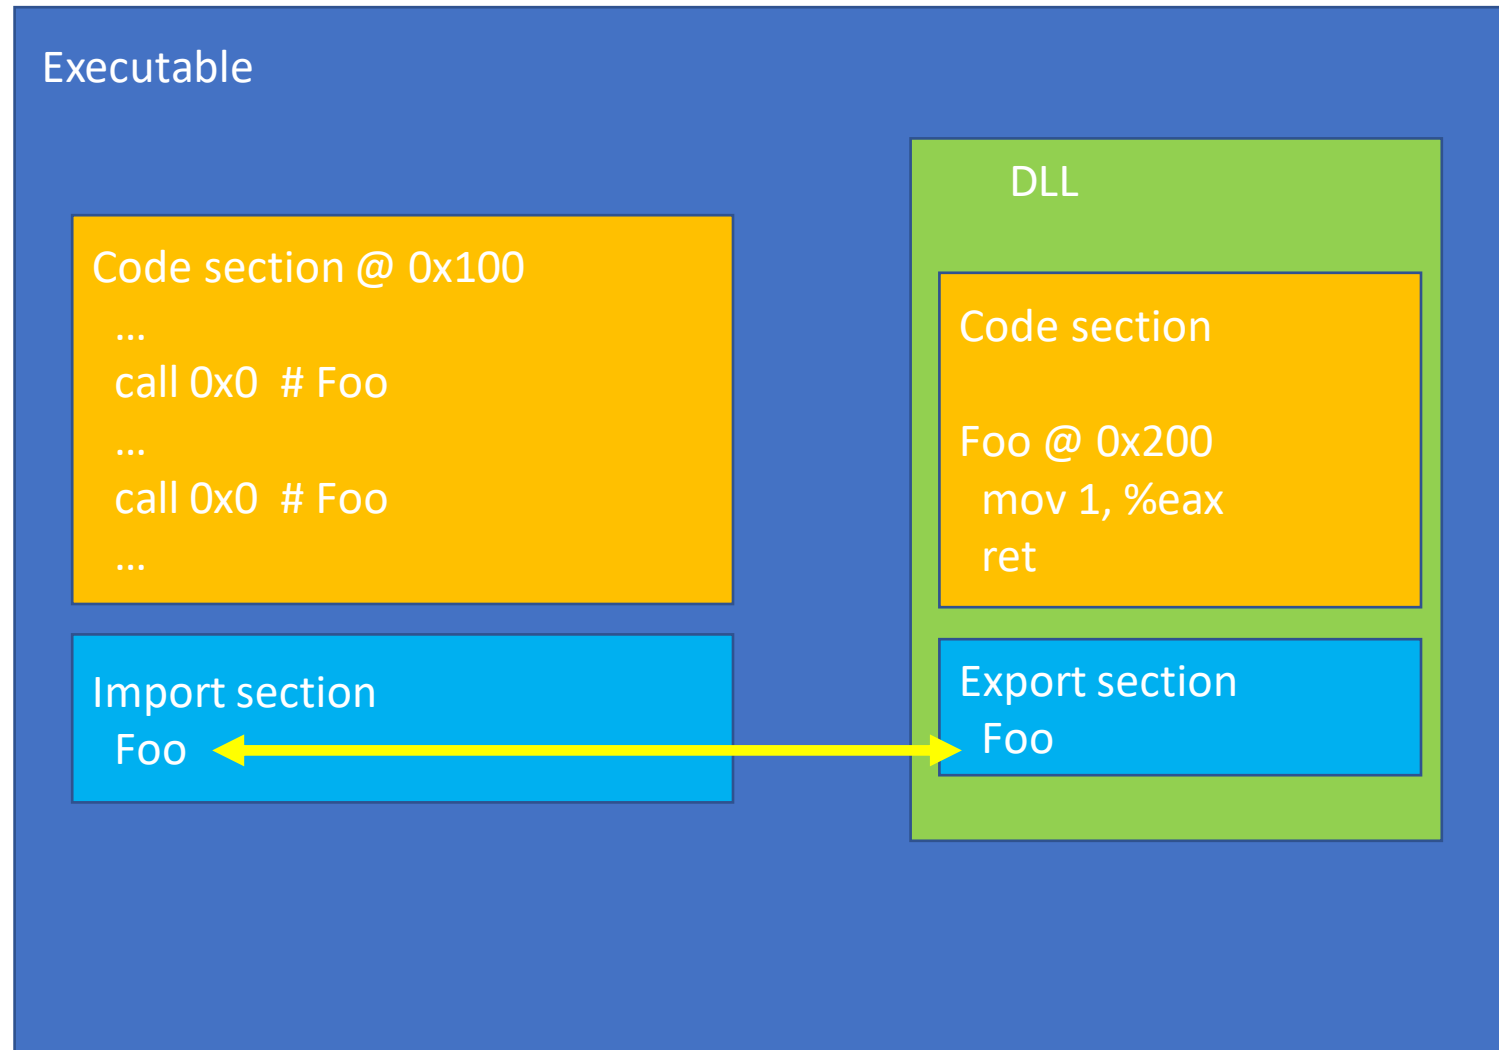
# Main principles

- DLLs and executables share the same format
  - PE on Windows, ELF on Linux
- Library keeps its exported symbols in a special table
  - .edata on Windows, .dynsym on Linux
- Executable file keeps the list of needed libraries and imported symbols in another table
  - .idata on Windows, .dynsym/.dynamic on Linux

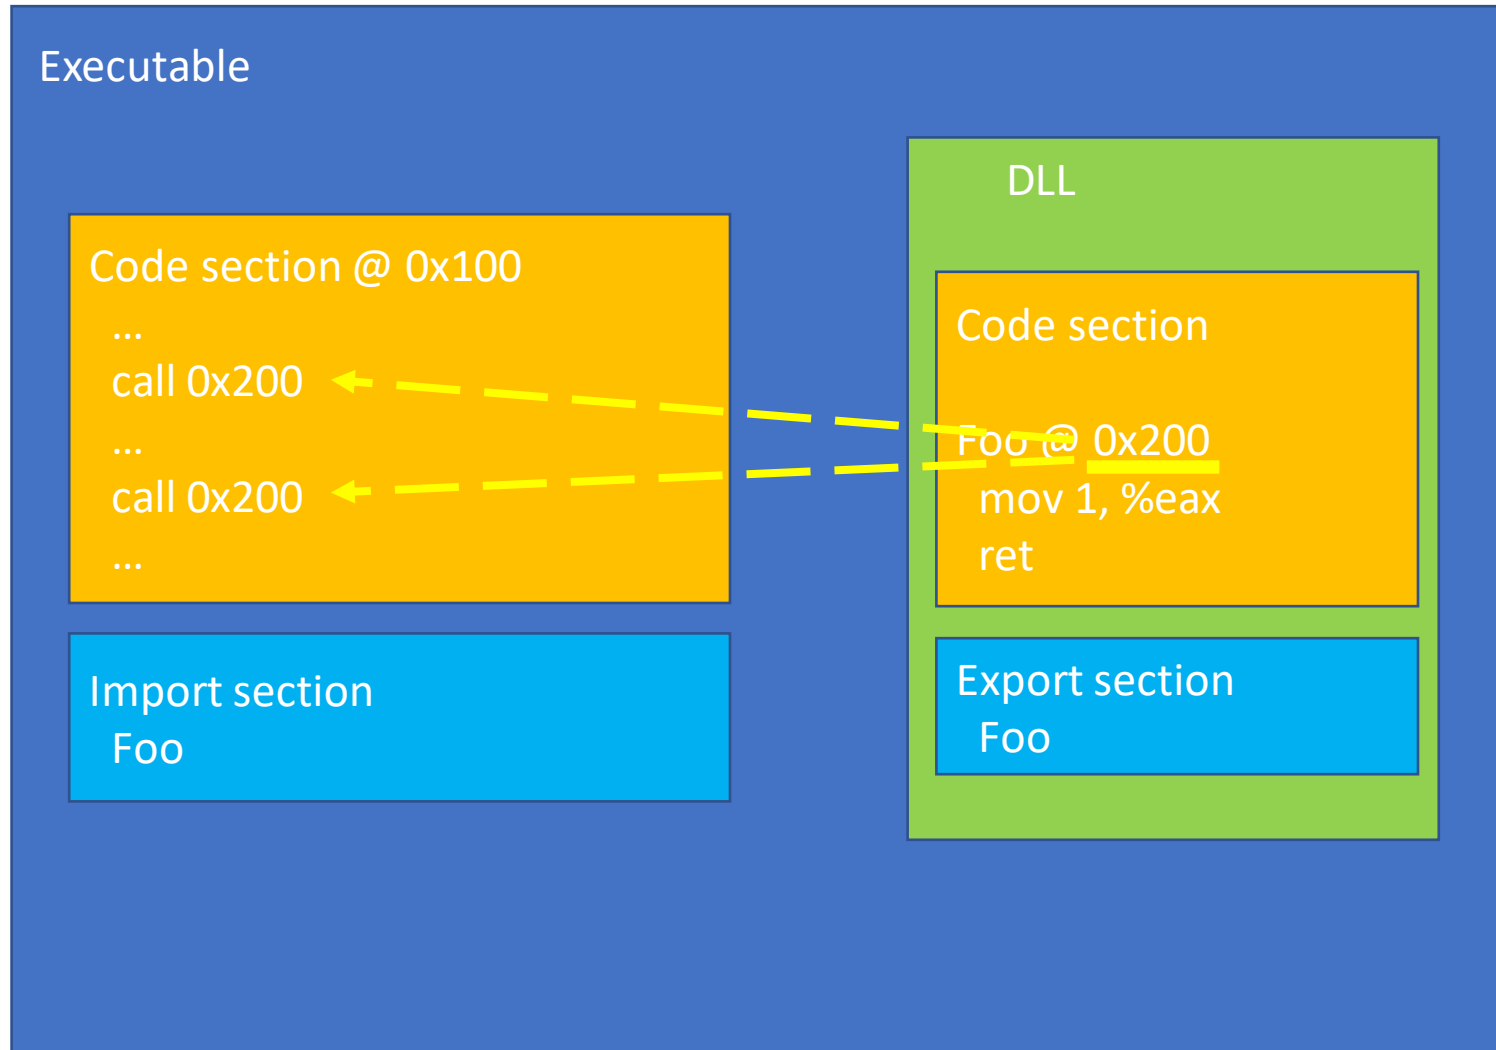
# Main principles



# Main principles

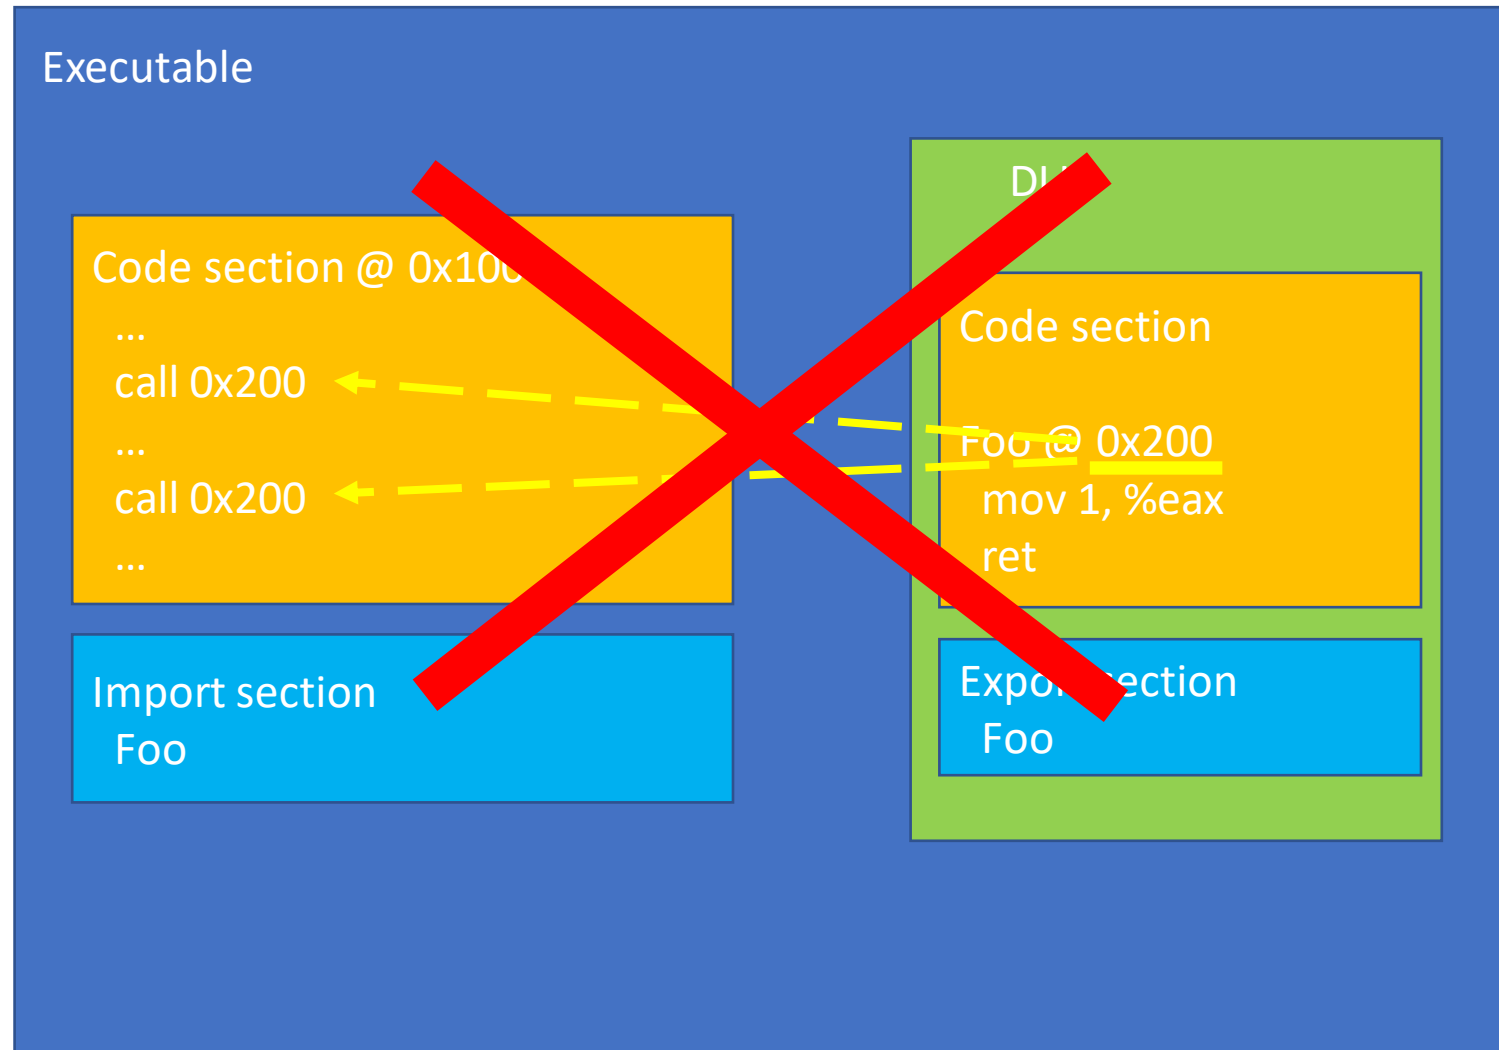


# Main principles

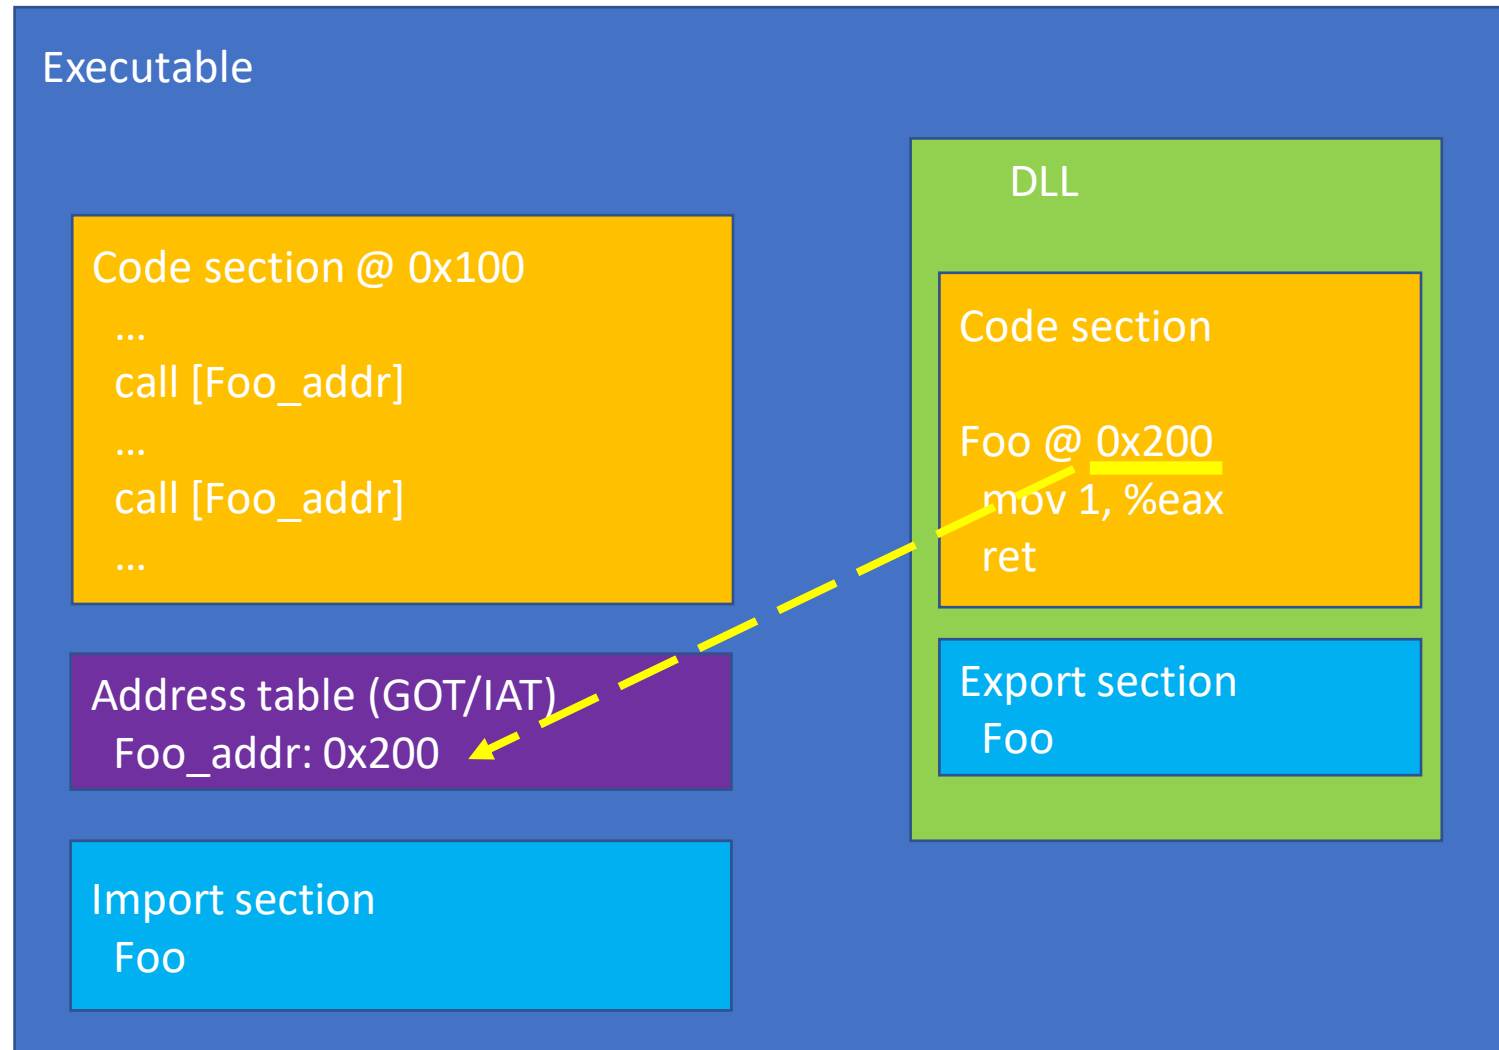




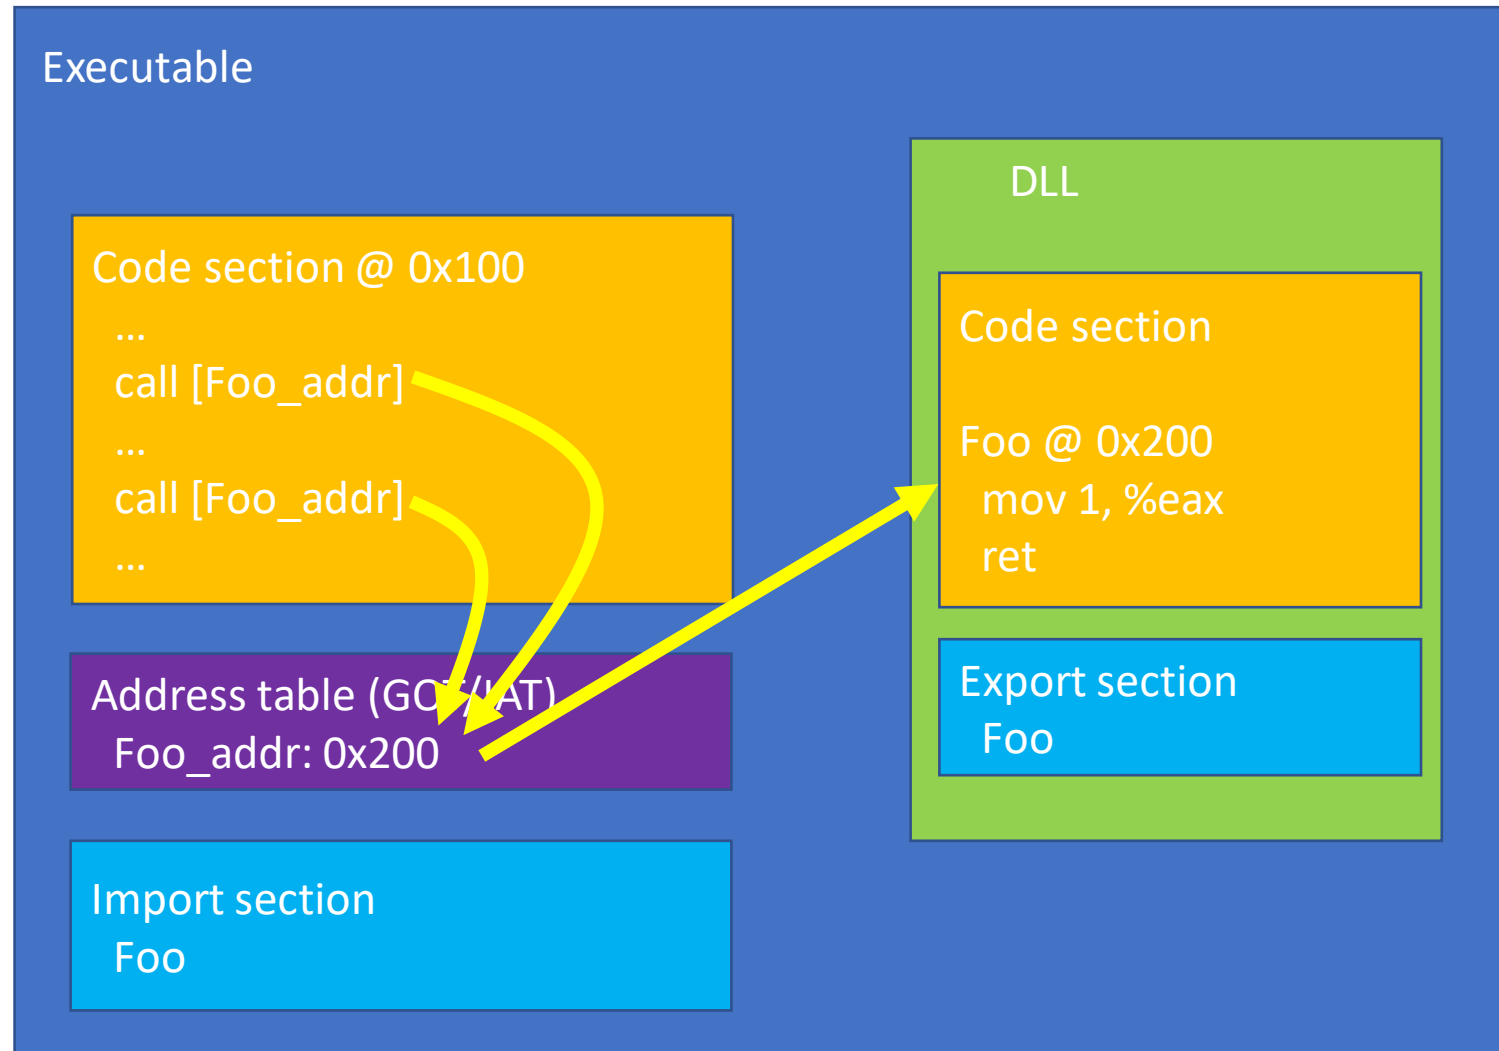
# Main principles



# Main principles




# Main principles



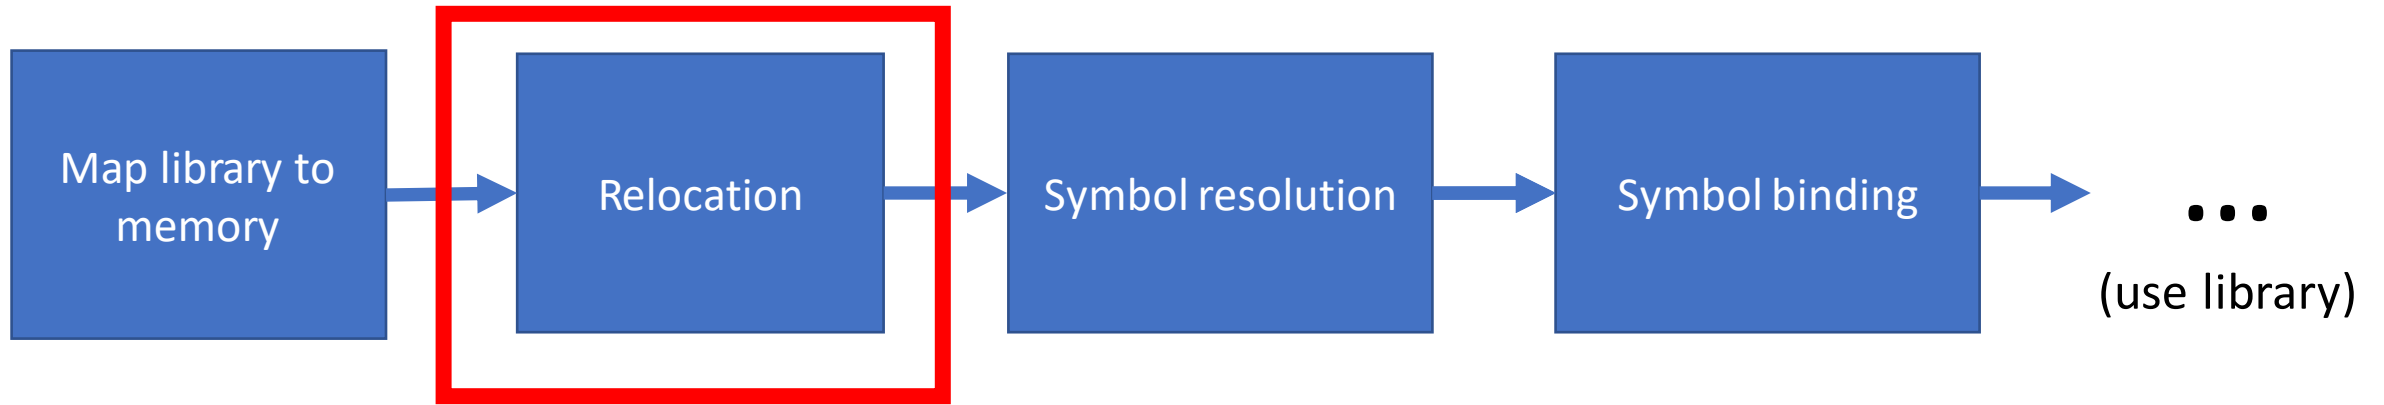
# Dynamic loader

- Libraries are imported into running program by dynamic loader
  - `/lib64/ld-linux-x86-64.so.2` on Linux
  - Image loader (Ldr) on Windows
- On program startup kernel maps loader into process memory and transfers control to it
- The loader
  - Maps needed libraries in process address space
  - Resolves and binds exported and imported symbols
  - Transfers control to main program

# Loading DLL



# Loading DLL



# Relocation: example

```
$ cat lib.c
```

```
int x = 0x12;
```

```
int *p = &x;
```

```
$ gcc -shared -fPIC lib.c
```

# Relocation: example

```
$ readelf --dyn-syms a.out
```

```
...
```

5:	000000000000004028	8	OBJECT	GLOBAL	DEFAULT	17	p
6:	000000000000004020	4	OBJECT	GLOBAL	DEFAULT	17	x



# Relocation: example

```
$ readelf --dyn-syms a.out
```

...

5:	00000000000004028	8	OBJECT	GLOBAL	DEFAULT	17	p
6:	00000000000004020	4	OBJECT	GLOBAL	DEFAULT	17	x

```
$ objdump -s -j .data a.out
```

4018 18400000 00000000 12000000 00000000 .@.....

4028 00000000 00000000 .....  


# Relocation

- Function and global variables addresses can only be determined at runtime
  - When library load address is known
- DLL contains special table with addresses of pointers that need to be updated after load
  - .rela.dyn on Linux, .reloc on Windows
- This patching is called *relocation*

# Relocation: example

```
$ readelf -r a.out
```

```
Relocation section '.rela.dyn' at offset 0x358 contains 8 entries:
```

```
...
```

```
000000004028 000600000001 R_X86_64_64 00000000000004020 x + 0
```



# Relocation: position-independent code


- All libraries are linked in position-independent (RIP/PC-relative) mode
  - Binary code does not use explicit addresses of function or global variables
  - Addresses are specified as offsets from to current instruction's address:

`mov global_var, %rdi`            `mov global_var(%rip), %rdi`

- Such code does not need to be relocated at load time
  - Faster library loads
  - Code segment is constant so can be shared by multiple processes
- Data still needs to be relocated (e.g. vtables)
  - But such relocations are few

```
int data;  
int *ptr = &data;    // Relocation needed
```

# Loading DLL



# Symbol resolution

- Matching exported and imported symbols
- To speed up search symbol information is stored in special hash tables
- Windows and Linux use different approaches:
  - Windows: each imported symbol is bound to particular library at link time (and will only be searched in that library)
  - Linux: imported symbols are searched sequentially in all loaded libraries
    - This enable runtime symbol interposition

# Runtime interposition

- We can force loader to find imported symbols in different library than the one it was supposed to come from
- Usually interposition is enabled via LD\_PRELOAD environment variable:

```
$ cat prog.c
int main(int argc) { printf("%d\n", argc); }
$ ./prog a b c
4
$ cat lib.c
int printf(char *fmt, ...) { puts("Hello from interceptor\n"); }
$ LD_PRELOAD=./lib.so ./prog a b c
Hello from interceptor
```

- Often used by debug tools like Electric Fence or AddressSanitizer to intercept memory operations (malloc, etc.)

# Interposition may hurt optimizations


- Compiler has to limit optimizations due to potential interposition
- E.g. compiler fails to inline due to potential interposition of foo:

```
$ cat mylib.c
void foo() {}
void bar() { foo(); }
```

```
$ gcc mylib.c -O3 -fPIC -S -o -
...
bar:
    jmp     foo@PLT
```



# Loading DLL



# Symbol binding

- Binding function calls in program with addresses of imported functions that were identified at symbol resolution stage
- Addresses of imported functions are stored in special dispatch table
  - Import Address Table on Windows, Global Offset Table on Linux
  - Initialized by loader at program startup
- Call of imported function is done by loading its address from the table:

```
# Windows
```

```
call qword ptr [__imp_foo]
```

```
# Linux
```

```
call *foo@GOTPCREL(%rip)
```

- Library calls are indirect (like virtual functions)

# Lazy binding on Linux

- Loading function address from dispatch table is done by a special stub function (PLT stub)
- PLT stubs are generated by linker
- Delays symbol resolution and binding until first use:

```
.section .text
...
call foo
...
.section .plt
foo:    # PLT stub pseudocode
if (first call)
    GOT[foo] = resolve address of foo
call GOT[foo]
```

Speeding up dynamic libraries

# DLL overheads

- Library load
  - Relocation
  - Symbol resolution and binding
- Library use
  - Indirect calls

# DLL overheads

- **Library load**
  - Relocation
  - Symbol resolution and binding
- Library use
  - Indirect calls

# DLL speedup: disabling unused libraries

- Often large programs may accidentally link against unused libraries
- Such libraries will slow program even down even if none of their functions are called
- `-Wl,--as-needed` flag allows linker to identify and ignore such libraries
- Enabled by default in some distros (Ubuntu but not Fedora/RHEL)

# DLL speedup: delayed library loading (lazy loading)

- Library may be used in only some rare scenarios
- Instead of loading it at startup we could load it on first use (lazy loading)
- Some platforms support this out-of-the-box:
  - Windows: /DELAYLOAD flag
  - macOS: -Wl,-z,-lazy-l flag (no longer supported)
- No standard solution for Linux but can use Implib.so
  - <https://github.com/yugr/Implib.so>





# Implib.so

- Given a DLL, generates small static library with stub functions (trampolines)
- Instead of DLL we link our program against that library
- At runtime executing a stub function will cause library to be loaded and control passed to it:

```
int foo(type1 arg1, type2 arg2, ...) { # Stub
    static void *foo_real = NULL;
    if (!foo_real) {
        void *handle = dlopen(...);
        foo_real = dlsym(handle, "foo");
    }
    return foo_real(arg1, arg2, ...);
}
```

# Implib.so

- Implements delayed loading for POSIX systems
- Uses runtime loading API (dlopen, dlsym)
- Supports many different targets
  - x86, ARM, AArch64, RISC-V, e2k, etc.
  - Linux (+ part. BSD)

# DLL overheads

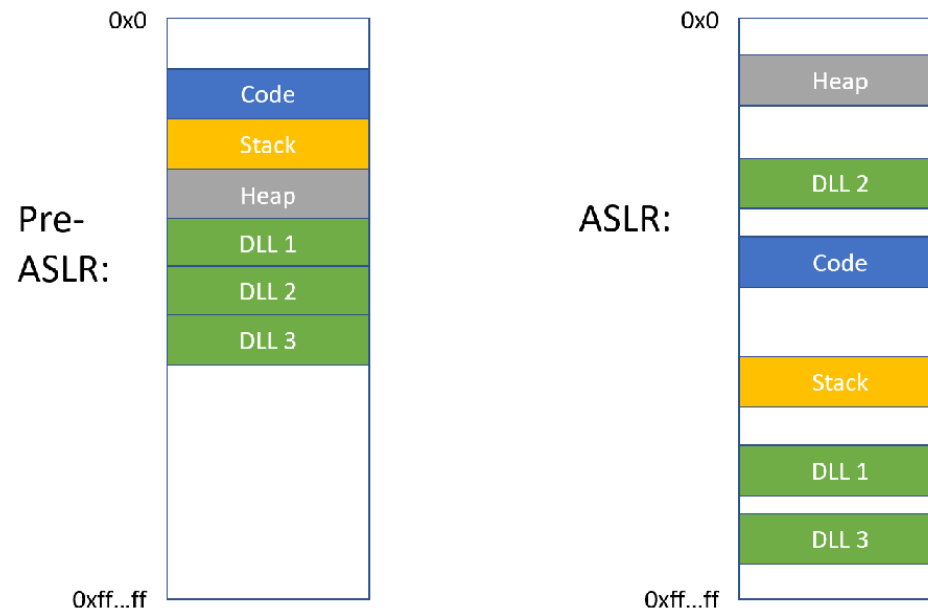
- Library load
  - **Relocation**
  - Symbol resolution and binding
- Library use
  - Indirect calls

# DLL speedup: link-time relocation

- Relocation may be avoided if library is linked at address which is guaranteed to be free at load time
- To achieve this we could
  - Scan all installed programs and libraries
  - Statically partition address space between all libraries
  - Link each library at its dedicated address
- Then dynamic loader will be able to avoid library relocation
- Solutions:
  - Windows: preferred load address (/BASE linker parameter)
  - Linux: Prelink

# DLL speedup: link-time relocation

- Optimization is no longer relevant due to modern security guidelines
- ASLR requires DLL to be loaded at random locations (to complicate hacker's work)



# Relocation: Windows optimization

- Legacy Win32-libraries do not use position-independent code
  - X86 lacks position-independent instructions
- A lot of instructions need to be relocated at load time
- To speed things up same library is loaded at the same address in all running processes
  - Relocation is only needed on first load
  - Works in modern Windows versions

# DLL overheads

- Library load
  - Relocation
  - **Symbol resolution and binding**
- Library use
  - Indirect calls

# DLL speedup: prelinking

- Dispatch table inside executable file could be statically initialized with precomputed function addresses
- This will only work if library is guaranteed to always be loaded at same fixed address
  - I.e. link-time relocation was performed
- Solutions:
  - Windows: DLL binding
  - Linux: Prelink
- No longer relevant in modern Windows and Linux due to ASLR



# DLL speedup: optimizing symbol tables

- During symbol resolution symbols are looked up in hashtables inside dynamic libraries
- On Linux linkers provide some means to control size and format of these hashtables
- The usually recommended set of options:
  - `-Wl,--hash-style=both -Wl,-O1`
- `-Wl,--hash-style=both` is turned on by default in all modern distros
- `-Wl,-O1` does not improve performance in practice

# DLL speedup: disable lazy binding

- Lazy binding on Linux speeds up program startup at the cost of additional function call overhead
- In addition to address load and indirect call we have a PLT stub call:
  - Extra jump
  - Increased cache/branch predictor pressure
  - Address has to be loaded from GOT on each call
- Lazy binding and related overheads may be disabled via `-fno-plt` compiler flag

# DLL speedup: disable dynamic loading

- -fno-plt
  - Speeds up library function calls
  - Reduces pressure on I\$ and BTB
  - Slows down program startup (as all addresses now need to be resolved and bound at program startup)
- Modern security guidelines suggest that all addresses are resolved at startup anyway
  - Allows Full Relro (-Wl,-z,relro) protection to avoid unintended GOT modifications during program run
  - Full Relro is enabled by default in RHEL/Fedora and Ubuntu

# DLL speedup: disable lazy binding

- Examples:
  - Using `-fno-plt` in Clang improves performance by up to 10%

# DLL overheads

- Library load
  - Relocation
  - Symbol resolution and binding
- Library use
  - **Indirect calls**

# Problem with exported symbols

- By default all library functions are exported on Linux
  - For compatibility with static libraries
- Due to potential interposition all function calls inside the library must go through GOT
- Overheads:
  - Unnecessary indirect calls
  - Disabled compiler optimizations (inlining, cloning, etc.)

# DLL speedup: disabling function interposition

- Special compiler flags allow compiler to ignore interposition
- `-Bsymbolic/-Bsymbolic-functions` – replaces indirect calls of library functions inside the library with direct calls
  - Turned on by default in some distributions (Ubuntu but not Debian)
- `-fno-semantic-interposition` – tells compiler to ignore possibility of interposition
  - Turned on by default in Clang but not GCC
  - Enabled in GCC under `-Ofast`
- Need both flags for optimal performance

# DLL speedup: disabling function interposition

- Examples:
  - Using -Bsymbolic-functions speeds up Clang by up to 10%
  - Using -fno-semantic-interposition when building Python gives up to 30% performance improvement
    - <https://fedoraproject.org/wiki/Changes/PythonNoSemanticInterpositionSpeedup>





# DLL speedup: reducing library interface

- Simple way to improve performance
- Does not require non-standard build flags
- Explicit control over which symbols are exported:

```
$ cat mylib.c  
void internal() {}
```

```
__attribute__((visibility("default")))  
void public() { internal(); }
```

```
$ gcc mylib.c -fvisibility=hidden -fPIC -shared
```

- Which functions to export?
  - Usually functions from public header files
  - Such functions are a tiny fraction of all library functions

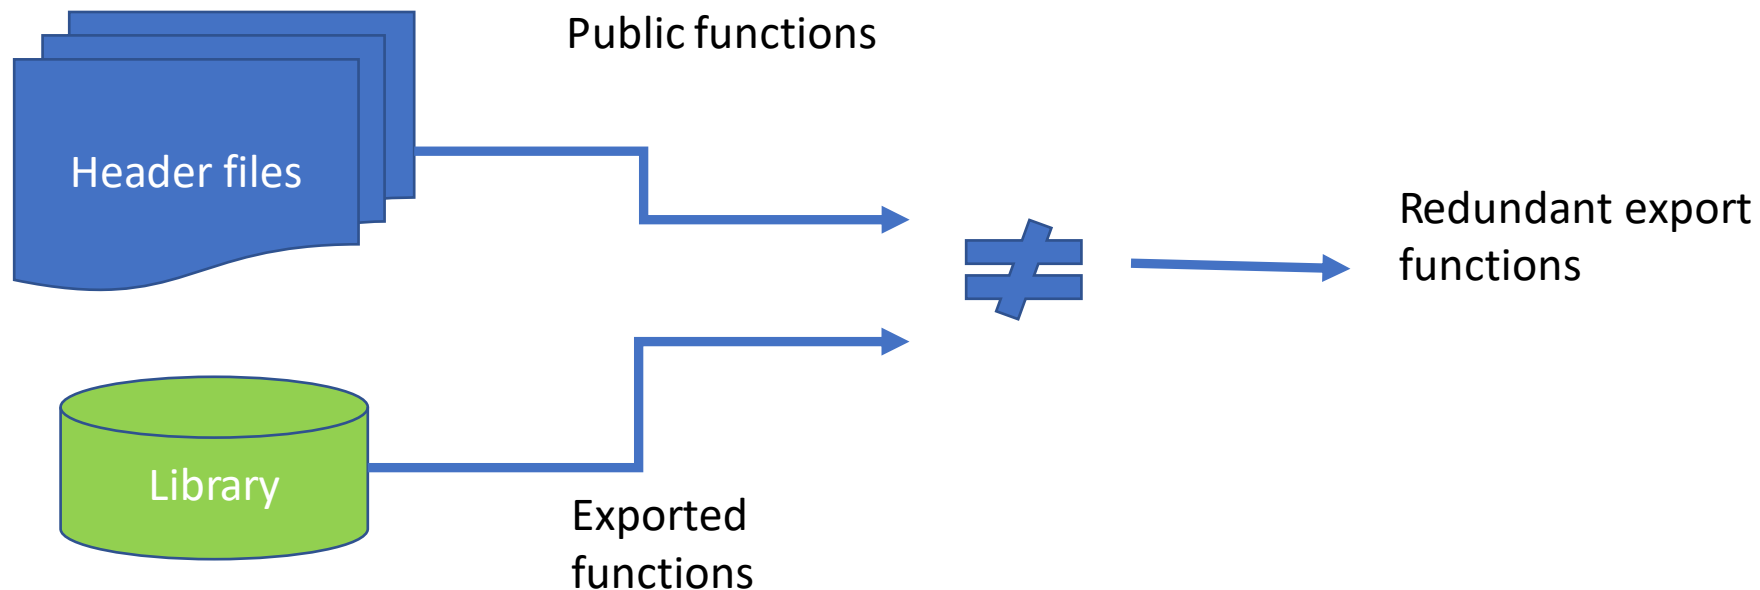
# DLL speedup: reducing library interface

- For a large code base (e.g. Linux distro) it may be hard to identify libraries with redundant exports
- Search of such libraries may be automated with ShlibVisibilityChecker tool
  - <https://github.com/yugr/ShlibVisibilityChecker>



# ShlibVisibilityChecker

- Analyzes functions in public library header files via libclang
- Compares them against functions actually exported by the library
- Reports redundant exports which need to be hidden



# ShlibVisibilityChecker example

```
$ read_header api --only-args /usr/include/x86_64-linux-  
gnu/gmp.h > api.txt
```

```
$ read_binary api --permissive /usr/lib/x86_64-linux-  
gnu/libgmp.so.10.4.1 > abi.txt
```

```
$ diff api.txt abi.txt | wc -l  
323
```

```
$ diff api.txt abi.txt
```

```
0a1,10
```

```
> __gmp_0
```

```
> __gmp_allocate_func
```

```
> __gmp_asprintf_final
```

```
> __gmp_asprintf_funs
```

```
...
```

# Conclusions

# Conclusions

- Dynamic libraries have some advantages over static ones

# Conclusions

- Dynamic libraries have some advantages over static ones
- But add overheads at program startup and program runtime

# Conclusions

- Dynamic libraries have some advantages over static ones
- But add overheads at program startup and program runtime
- Modern toolchains provide means to reduce overheads
  - Especially on Linux



# Additional reading

- Linkers, Loaders and Shared Libraries in Windows, Linux, and C++ (Ofek Shilon, CppCon 2023)
  - <https://www.youtube.com/watch?v=enXulxuNV4>
  - General overview of DLLs on different platforms
- How to Write Shared Libraries (by Ulrich Drepper)
  - <https://www.akkadia.org/drepper/dsohowto.pdf>
  - All you need to know about DLLs on Linux
- Everything You Ever Wanted to Know about DLLs (by James McNellis, CppCon 2017)
  - <https://www.youtube.com/watch?v=JPQWQfDhICA>
  - All you need to know about DLLs on Windows
- MaskRay Blog
  - <https://maskray.me/blog>
  - Linux system programming blog (GOT, PLT, etc.)



Thanks!

# Check RAM savings

- Build scanner
  - `gcc -Wall -Wextra scripts/ram-savings.c`
- Run under sudo:
  - `sudo ./a.out`

# Check disk savings

- Run
  - `scripts/disk-savings.pl`
- Script reports upper bound – real savings would be lower
  - With static libs not all functions will be imported by the applications
  - So only parts of the libraries will be included in executables

# Check -Wl,-O1

- Build two versions of LLVM:
  - -DBUILD\_SHARED\_LIBS=ON
  - -DBUILD\_SHARED\_LIBS=ON -DCMAKE\_SHARED\_LINKER\_FLAGS='-Wl,-O1'
- Compare performance:
  - ./benchmark.pl 10 path/to/clang -h

# Check -fno-plt

- Build two versions of LLVM:
  - -DBUILD\_SHARED\_LIBS=ON
  - -DBUILD\_SHARED\_LIBS=ON -DCMAKE\_CXX\_FLAGS='-fno-plt'
- Compare performance:
  - ./benchmark.pl 10 path/to/clang -S -O2 ~/InstCombining.ii

# Check -Bsymbolic-functions

- Build two versions of LLVM:
  - -DBUILD\_SHARED\_LIBS=ON
  - -DBUILD\_SHARED\_LIBS=ON -DCMAKE\_SHARED\_LINKER\_FLAGS='-Wl,-Bsymbolic-functions'
- Compare performance:
  - `./benchmark.pl 10 path/to/clang -S -O2 ~/InstCombining.ii`

# Address-space Layout Randomization

